

# ATE Estimation using Full Propensity Matching (implemented from scratch) and Linear Propensity Scoring

Amir Idris

```
if(!require("dplyr")){
  install.packages("dplyr")
}

## Loading required package: dplyr

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

if(!require("MatchIt")){
  install.packages("MatchIt")
}

## Loading required package: MatchIt

if(!require("lmtest")){
  install.packages("lmtest")
}

## Loading required package: lmtest
## Loading required package: zoo

##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric

if(!require("sandwich")){
  install.packages("sandwich")
}

## Loading required package: sandwich

if(!require("boot")){
  install.packages("boot")
}
```

```

## Loading required package: boot
if(!require("survival")){
  install.packages("survival")
}

## Loading required package: survival
##
## Attaching package: 'survival'
## The following object is masked from 'package:boot':
##
##      aml
if(!require("optmatch")){
  install.packages("optmatch")
}

## Loading required package: optmatch
## The optmatch package has an academic license. Enter relaxinfo() for more information.
if(!require("glmnet")){
  install.packages("glmnet")
}

## Loading required package: glmnet
## Loading required package: Matrix
## Loaded glmnet 4.1-1
if(!require("gbm")){
  install.packages("gbm")
}

## Loading required package: gbm
## Loaded gbm 2.1.8
if(!require("rpart")){
  install.packages("rpart")
}

## Loading required package: rpart
library(dplyr)
library(MatchIt)
library(lmtest)
library(sandwich)
library(boot)
library(survival)
library(optmatch)
library(glmnet)
library(gbm)
library(rpart)

```

First, let's load the data and summarize it.

```

lowdim_data <- read.csv("../data/lowDim_dataset.csv")
highdim_data <- read.csv("../data/highDim_dataset.csv")

```

```
print(dim(lowdim_data))
```

```
## [1] 500 24
```

```
print(dim(highdim_data))
```

```
## [1] 2000 187
```

We can see that the high-dimensional data has an order of magnitude more dimensions, and four times the data as compared to the low dimensional data. Let's view a couple rows

```
head(lowdim_data)
```

```
##           Y A  V1  V2  V3  V4 V5  V6 V7  V8  V9 V10 V11 V12 V13 V14
## 1 30.48700 0 0.00 0.00 0.00 0.0 0 0.00 0.0 0.00 0.00 0.00 0.00 0 0.00 0.00
## 2 18.20842 0 0.00 0.00 0.00 0.0 0 0.00 0.0 1.40 0.00 0.00 0.00 0 0.70 0.00
## 3 13.48504 0 0.00 0.00 0.00 0.0 0 0.00 0.0 0.00 0.00 0.00 0.00 0 0.00 0.00
## 4 25.69968 1 2.38 0.00 0.00 0.0 0 0.00 0.0 0.00 0.00 0.00 0.00 0 0.00 0.00
## 5 23.75297 0 0.15 0.15 0.05 0.1 0 0.42 0.1 0.95 0.42 0.05 0.05 0 0.00 0.36
## 6 13.63108 0 0.16 0.00 0.00 0.0 0 0.16 0.0 0.16 0.16 0.00 0.00 0 0.16 0.00
##      V15 V16 V17 V18 V19 V20 V21      V22
## 1 0.00 0 0.00 0 0.00 0.00 9.09 1.1496222
## 2 1.40 0 1.40 0 0.00 0.00 0.00 2.8877015
## 3 3.57 0 0.00 0 0.00 0.00 0.00 0.0000000
## 4 2.38 0 2.38 0 0.00 0.00 0.00 0.4054651
## 5 3.16 0 1.58 0 0.52 0.31 0.00 1.5746394
## 6 2.88 0 0.50 0 0.00 0.00 0.00 0.3798054
```

```
head(highdim_data)
```

```
##           Y A V1 V2 V3 V4 V5 V6 V7      V8 V9 V10 V11 V12 V13 V14 V15 V16 V17 V18
## 1 41.224513 0 0 1 4 18 17 -1 1 0.75 1 28 2 20 20 15 15 5 25 1
## 2 40.513875 0 0 0 1 10 6 -1 10 0.35 1 30 2 35 15 25 10 5 10 1
## 3 38.495476 0 0 0 16 8 4 4 4 0.40 1 26 2 30 25 20 10 10 5 1
## 4 33.001889 0 1 0 3 10 2 -1 5 0.41 0 24 4 5 15 45 25 0 10 1
## 5 37.043603 0 1 1 11 21 10 10 20 0.43 1 28 2 25 20 20 25 0 10 1
## 6 -2.877098 1 0 1 11 21 14 14 15 -0.20 1 28 1 20 18 20 17 10 15 1
##      V19 V20 V21 V22 V23 V24 V25 V26 V27 V28 V29 V30 V31 V32      V33 V34 V35 V36 V37
## 1 9 -1 8 8 7 5 7 6 2 13 141 10 1 1 0 2 1 3 205
## 2 7 6 3 5 6 5 6 4 2 7 152 2 1 1 0 2 1 2 234
## 3 8 9 9 7 9 7 7 8 2 12 239 12 9 29 25000 2 0 2 66
## 4 7 7 8 5 5 5 6 3 2 11 189 7 1 1 0 1 1 2 226
## 5 6 8 8 7 5 -1 8 5 2 11 242 4 54 59 25552 2 3 6 262
## 6 6 10 10 8 7 6 7 7 2 10 223 4 240 64 26100 1 9 11 40
##      V38 V39 V40 V41 V42 V43 V44 V45 V46 V47 V48 V49 V50 V51 V52 V53 V54 V55 V56
## 1 17 1 211 7 7 2 8 10 9 9 9 2 9 10 4 10 10 8 10
## 2 103 1 188 1 1 1 6 7 6 7 7 5 7 7 7 9 7 8 7
## 3 50 1 315 11 6 2 5 8 5 6 5 2 9 9 8 9 9 6 6
## 4 198 3 225 6 3 2 2 7 8 8 3 3 7 10 7 5 9 7 7
## 5 195 3 258 2 9 5 3 8 5 5 6 4 4 9 2 6 7 5 7
## 6 11 2 353 10 1 1 1 5 7 7 7 1 3 9 4 7 7 7 8
##      V57 V58 V59 V60 V61 V62 V63 V64 V65 V66 V67 V68 V69 V70 V71 V72 V73 V74 V75
## 1 6 9 5 6 10 20 20 20 20 10 -1 -1 -1 -1 -1 -1 30 10 10
## 2 1 8 1 2 20 20 20 20 10 10 -1 -1 -1 -1 -1 -1 30 10 15
## 3 4 7 5 -1 5 30 20 10 5 30 50 5 10 20 5 10 50 10 5
## 4 7 1 5 2 21 17 22 20 8 13 -1 -1 -1 -1 -1 -1 18 15 17
```

```

## 5  2  2  5 -1 20 20 15 25 10 10 40 5 10 20 5 20 30 20 5
## 6  9  6  1 -1 20 20 20 10 20 10 10 30 30 5 20 5 50 30 5
##   V76 V77 V78 V79 V80 V81 V82 V83 V84 V85 V86 V87 V88 V89 V90 V91 V92 V93 V94
## 1  30 10 10  9 10  9  9  9 -1 -1 -1 -1 -1  6  9  9  8  8  5
## 2  30  5 10  7  8  9  7  8 -1 -1 -1 -1 -1  7 10 10  7 10  5
## 3  10  5 20  9  8  5  9  3  8  8  9  3  8  4  7  6  6  5  8
## 4  14 20 16  6  8  7  8  8 -1 -1 -1 -1 -1  4  7  7  6  5  3
## 5  25 10 10  5  8  7  8  7  5  8  7  6  6  7  6  7  6  5  5
## 6  5  5  5  8  8  8  8  5  8  7 10  9  8 10  8  9 10 10  7
##   V95 V96 V97 V98 V99 V100 V101 V102 V103 V104 V105 V106 V107 V108 V109 V110
## 1  7  8  2  3 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1  6
## 2  5  1  2  2 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1  4
## 3  6  8  2  3 20 20 20 10  5 25  8  8  8  4  4 -1
## 4  6  5  2 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1  4
## 5  6  3  0  1 35 15 20 20  5  5  4  7  5  4  7  4
## 6  9  4  0  2 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1  5
##   V111 V112 V113 V114 V115 V116 V117 V118 V119 V120 V121 V122 V123 V124
## 1  1  2 -1 -1 -1 -1 -1 -1 20.00 20.00 15.00 20.00 10.00 15.00
## 2  3  2 -1 -1 -1 -1 -1 -1 24.14 13.79 20.69 27.59 10.34 3.45
## 3 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1.00 -1.00 -1.00 -1.00 -1.00
## 4  1  2 -1 -1 -1 -1 -1 -1 25.00 10.00 20.00 20.00  5.00 20.00
## 5  1  2 -1 -1 -1 -1 -1 -1 30.00 20.00 15.00 30.00  5.00  0.00
## 6  3  2 -1 -1 -1 -1 -1 -1 20.00 20.00 20.00 20.00 10.00 10.00
##   V125 V126 V127 V128 V129 V130 V131 V132 V133 V134 V135 V136 V137 V138
## 1 -1 -1 -1 -1 -1 -1 -1.00 -1.00 -1.00 -1.00 -1.00 -1.00  7  9
## 2 -1 -1 -1 -1 -1 -1 -1.00 -1.00 -1.00 -1.00 -1.00 -1.00  6  8
## 3 -1 -1 -1 -1 -1 -1 -1.00 -1.00 -1.00 -1.00 -1.00 -1.00 -1 -1
## 4 -1 -1 -1 -1 -1 -1 -1.00 -1.00 -1.00 -1.00 -1.00 -1.00  6  7
## 5 40 15 15 30  0  0 33.33 11.11  5.56 33.33 16.67  0.00  5  7
## 6 10 30 10 10 30 10 36.36 18.18  9.09 18.18  4.55 13.64  8  8
##   V139 V140 V141 V142 V143 V144 V145 V146 V147 V148 V149 V150 V151 V152 V153
## 1  9  9  8 -1 -1 -1 -1 -1  1  4  1 -1 -1 15 20
## 2  7  8  6 -1 -1 -1 -1 -1  0  0  0 -1 -1 20 20
## 3 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
## 4  8  9  8 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
## 5  7  5  7  5  7  7  4  7  0  4  1  2 -1 30 20
## 6  8 10  6  8  8  9 10  8 -1 -1 -1 -1 -1 -1 -1
##   V154 V155 V156 V157 V158 V159 V160 V161 V162 V163 V164 V165 V166 V167 V168
## 1 15 20 15 15 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
## 2 20 20  0 20 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
## 3 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
## 4 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
## 5 20 20  5  5 30 20 20 30  0  0 40 10 20 20  5
## 6 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
##   V169 V170 V171 V172 V173 V174 V175 V176 V177 V178 V179 V180 V181 V182 V183
## 1 -1 -1 -1 -1 -1 -1 -1  8 10  8  9  8 -1 -1 -1
## 2 -1 -1 -1 -1 -1 -1 -1  6  5  6  8  5 -1 -1 -1
## 3 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
## 4 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
## 5  5 30 10 20 10 20 10  6  8  7  7  7  6  7  6
## 6 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1
##   V184 V185
## 1 -1 -1
## 2 -1 -1

```

```
## 3   -1   -1
## 4   -1   -1
## 5    6    7
## 6   -1   -1
```

First, we must train models to estimate the propensity scores. We have five models assigned to us:

P1: Logistic Regression  
P2: L1 Penalized Logistic Regression  
P3: L2 penalized logistic regression  
P4: Regression trees  
P5: Boosted stumps

```
# Model selector method
source("../lib/LR_propensity_est.R")
source("../lib/boosted_stumps_propensity_est.R")
source("../lib/regression_tree_propensity_est.R")
model_selector <- function(data, mode = 1){
  if (mode == 1) {
    return(lr_propensity_model(data))
  } else if (mode == 2 | mode == 6){
    return(lr_propensity_model(data, mode = "lasso"))
  } else if (mode == 3){
    return(lr_propensity_model(data, mode = "ridge"))
  } else if (mode == 4){
    return(regression_tree_propensity_model(data))
  } else{
    return(gbm_propensity_model(data))
  }
}
```

```
# P1-P5, A7 + P2 included
m <- 6 # number of models
lowdim_models <- list()
highdim_models <- list()
lowdim_model_times <- list()
highdim_model_times <- list()
for(i in 1:m){
  lowdim_model_times[[i]] <- system.time({lowdim_models[[i]] <- model_selector(lowdim_data, mode = i)};})
}
for(i in 1:m){
  highdim_model_times[[i]] <- system.time({highdim_models[[i]] <- model_selector(highdim_data, mode = i)};})
}
```

Then, we'll estimate propensity scores, and transform them to use linear propensity distance.

*\*\*It's worth noting that we can't have propensity scores that are exactly 0 or exactly 1, since that will give us infinity for our logit transformation. So, the entries from the regression tree that match this criteria are added/subtracted a small dx so that they don't go to infinity.*

```
lowdim_prop_scores <- list()
highdim_prop_scores <- list()
lowdim_score_times <- list()
highdim_score_times <- list()

predict_selector <- function(features, model, mode = 1){
  if (mode == 1) {
```

```

    return(lr_propensity(features, model))
  } else if (mode == 2 | mode == 6){
    return(lr_propensity(features, model, mode = mode))
  } else if (mode == 3){
    return(lr_propensity(features, model, mode = mode))
  } else if (mode == 4){
    return(regression_tree_propensity(features, model))
  } else{
    return(gbm_propensity(features, model))
  }
}

for(i in 1:m){
  lowdim_score_times[[i]] <- system.time({
    if(i < 6){
      lowdim_prop_scores[[i]] <- logit(predict_selector(lowdim_data, lowdim_models[[i]], mode = i))
    }
    else{
      lowdim_prop_scores[[i]] <- predict_selector(lowdim_data, lowdim_models[[i]], mode = i)
    };})
}

for(i in 1:m){
  highdim_score_times[[i]] <- system.time({
    if (i < 6){
      highdim_prop_scores[[i]] <- logit(predict_selector(highdim_data, highdim_models[[i]], mode = i))
    }
    else{
      highdim_prop_scores[[i]] <- predict_selector(highdim_data, highdim_models[[i]], mode = i)
    };})
}

```

Now, we'll proceed with the full matching, using the *MatchIt* package.

This part is the most time intensive, since it must find the optimal matching such that total discrepancy between matched comparisons is minimized, so it can't take a greedy approach. (*Hansen, 2004*)

```

m <- 5 # We'll handle weighted regression separately
lowdim_data_match <- subset(lowdim_data, select = -c(Y))
highdim_data_match <- subset(highdim_data, select = -c(Y))
lowdim_matches <- list()
highdim_matches <- list()
lowdim_match_times <- list()
highdim_match_times <- list()

for(i in 1:m){
  lowdim_match_times[[i]] <- system.time({lowdim_matches[[i]] <- matchit(A ~ ., data = lowdim_data_match)})
}

for(i in 1:m){
  highdim_match_times[[i]] <- system.time({highdim_matches[[i]] <- matchit(A ~ ., data = highdim_data_match)})
}

# Example matchings
lowdim_matches[[1]]

## A matchit object

```

```
## - method: Optimal full matching
## - distance: User-defined
## - number of obs.: 500 (original), 500 (matched)
## - target estimand: ATE
## - covariates: V1, V2, V3, V4, V5, V6, V7, V8, V9, V10, V11, V12, V13, V14, V15, V16, V17, V18, V19,
```

```
highdim_matches[[1]]
```

```
## A matchit object
## - method: Optimal full matching
## - distance: User-defined
## - number of obs.: 2000 (original), 2000 (matched)
## - target estimand: ATE
## - covariates: too many to name
```

We obtain datasets of the matches, and bind our outcome Y back to them.

```
lowdim_match_sets <- list()
highdim_match_sets <- list()
for(i in 1:m){
  lowdim_match_sets[[i]] <- match.data(lowdim_matches[[i]])
  Y.low <- lowdim_data$Y
  lowdim_match_sets[[i]] <- as.data.frame(cbind(Y.low, lowdim_match_sets[[i]]))
}
for(i in 1:m){
  highdim_match_sets[[i]] <- match.data(highdim_matches[[i]])
  Y.high <- highdim_data$Y
  highdim_match_sets[[i]] <- as.data.frame(cbind(Y.high, highdim_match_sets[[i]]))
}
```

Finally, we estimate our ATEs.

```
lowdim_ATEs <- c()
highdim_ATEs <- c()

for(i in 1:m){
  fit.low <- lm(Y.low ~ A, data = lowdim_match_sets[[i]], weights = weights)
  coeftest(fit.low, vcov. = vcovCL, cluster = ~subclass)
  lowdim_ATEs <- c(lowdim_ATEs, summary(fit.low)$coefficients["A", "Estimate"])
}
for(i in 1:m){
  fit.high <- lm(Y.high ~ A, data = highdim_match_sets[[i]], weights = weights)
  coeftest(fit.high, vcov. = vcovCL, cluster = ~subclass)
  highdim_ATEs <- c(highdim_ATEs, summary(fit.high)$coefficients["A", "Estimate"])
}
```

```
# Weighted regression
source("../lib/weighted_regression.R")
lowdim_match_times[[m+1]] <- system.time({lowdim_ATEs[[m+1]] <- weighted_regression(lowdim_data, lowdim_
highdim_match_times[[m+1]] <- system.time({highdim_ATEs[[m+1]] <- weighted_regression(highdim_data, high
```

How do our ATE estimates compare to the real ones?

```
m <- 6
# Real ATEs
lowdim_ATE_real <- 2.0901
highdim_ATE_real <- -54.8558
```

```

#SDs of outcomes, so we can normalize our difference in estimation
sd.low <- sd(Y.low)
sd.high <- sd(Y.high)

ATE_diff.low <- (lowdim_ATEs - lowdim_ATE_real)/sd.low
ATE_diff.high <- (highdim_ATEs - highdim_ATE_real)/sd.high

method_names <- c("A1 + D3 + P1", "A1 + D3 + P2", "A1 + D3 + P3", "A1 + D3 + P4", "A1 + D3 + P5", "A7 + P2")

ATE_table <- as.matrix(cbind(lowdim_ATEs, ATE_diff.low, highdim_ATEs, ATE_diff.high))
rownames(ATE_table) <- method_names
colnames(ATE_table) <- c("Low-Dim ATE Estimate", "Low-Dim Error in SDs", "High-Dim ATE Estimate", "High-Dim Error in SDs")
ATE_table

```

```

##           Low-Dim ATE Estimate Low-Dim Error in SDs High-Dim ATE Estimate
## A1 + D3 + P1           1.725128          -2.925137e-02          -60.14118
## A1 + D3 + P2           1.171943          -7.358735e-02          -57.18506
## A1 + D3 + P3           1.666441          -3.395494e-02          -59.20897
## A1 + D3 + P4           4.369954           1.827230e-01          -52.62841
## A1 + D3 + P5           2.371125           2.252324e-02          -53.38666
## A7 + P2                2.090175           6.028979e-06          -57.95906
##           High-Dim Error in SDs
## A1 + D3 + P1          -0.09748049
## A1 + D3 + P2          -0.04295943
## A1 + D3 + P3          -0.08028740
## A1 + D3 + P4           0.04108073
## A1 + D3 + P5           0.02709601
## A7 + P2              -0.05723466

```

Not bad! Our worst estimate is less than 0.2 standard deviations away from the true ATE.

Now, let's take a look at the time our algorithm took:

```

lowdim_model_time <- c()
highdim_model_time <- c()
lowdim_score_time <- c()
highdim_score_time <- c()
lowdim_match_time <- c()
highdim_match_time <- c()

for(i in 1:m){
  lowdim_model_time <- c(lowdim_model_time, lowdim_model_times[[i]][3])
  highdim_model_time <- c(highdim_model_time, highdim_model_times[[i]][3])
  lowdim_score_time <- c(lowdim_score_time, lowdim_score_times[[i]][3])
  highdim_score_time <- c(highdim_score_time, highdim_score_times[[i]][3])
  lowdim_match_time <- c(lowdim_match_time, lowdim_match_times[[i]][3])
  highdim_match_time <- c(highdim_match_time, highdim_match_times[[i]][3])
}

time_table <- as.matrix(cbind(lowdim_model_time, highdim_model_time, lowdim_score_time, highdim_score_time, lowdim_match_time, highdim_match_time))
total <- rowSums(time_table)
time_table <- as.matrix(cbind(time_table, total))
rownames(time_table) <- method_names
colnames(time_table) <- c("Low-Dim Model Train Time", "High-Dim Model Train Time", "Low-Dim Scoring Time", "High-Dim Scoring Time", "Low-Dim Match Time", "High-Dim Match Time")

```



time\_table

```
##          Low-Dim Model Train Time High-Dim Model Train Time
## A1 + D3 + P1          0.011          0.356
## A1 + D3 + P2          0.034          3.453
## A1 + D3 + P3          0.022          0.561
## A1 + D3 + P4          0.043          0.859
## A1 + D3 + P5          0.024          0.360
## A7 + P2          0.011          3.382
##          Low-Dim Scoring Time High-Dim Scoring Time Low-Dim Matching Time
## A1 + D3 + P1          0.003          0.005          0.160
## A1 + D3 + P2          0.034          0.018          0.153
## A1 + D3 + P3          0.004          0.014          0.158
## A1 + D3 + P4          0.022          0.007          0.136
## A1 + D3 + P5          0.003          0.027          0.147
## A7 + P2          0.004          0.013          0.063
##          High-Dim Matching Time Total
## A1 + D3 + P1          3.208 3.743
## A1 + D3 + P2          2.661 6.353
## A1 + D3 + P3          2.643 3.402
## A1 + D3 + P4          2.317 3.384
## A1 + D3 + P5          2.622 3.183
## A7 + P2          0.962 4.435
```

As we can see with quadruple the data, the full matching algorithm took about 16.781 times longer for the high-dimensional dataset, and the other components of the process were negligible in comparison. So, we can hypothesis that we have an  $O(n^2)$  time complexity.

Weighted Regression seems to be the overall champion, with a low runtime compared to full-matching, and very good estimates, with an almost exact estimate on the low-dimensional dataset.

References:

- <https://cran.r-project.org/web/packages/MatchIt/vignettes/estimating-effects.html#after-full-matching>
- Hansen BB. Full matching in an observational study of coaching for the SAT. Journal of the American Statistical Association. 2004;99(467):609–618.